

RAKSHIT NAIDU NEMAKALLU

[Email](#) ◇ [LinkedIn](#) ◇ [Website](#) ◇ [Google Scholar](#)

OBJECTIVE

My research interests hover around topics in Ethical Machine Learning (ML), Trustworthy Artificial Intelligence (AI), and AI for societal good specifically pertaining to solidarity. I'm interested in creating applications that have a direct impact on society through my research. Hence, this has led me to pursue understanding ML models through the following important questions : (1) How do we quantify disparities that arise in decisions made by ML models? (2) How do we create novel algorithms which mitigate the above effects?

EDUCATION

Master of Science (M.Sc.) in Information Technology (Privacy Engineering), Carnegie Mellon University
2021 - 2022

Selected Courses: [Ethics in Machine Learning](#), [Foundations of Privacy](#), Privacy Policy, Law and Technology (PPLT) and [ML with Large Datasets](#)

Bachelor of Technology (B.Tech.) in Computer Science and Engineering, Manipal Institute of Technology
2017 - 2021

Minor in Computational Mathematics

Selected Courses: Computational Linear Algebra, Distributed and Cloud Computing, Graph Theory and Matrices.

EXPERIENCE

Graduate Research Intern

Carnegie Mellon University

Apr 2023 - July 2023

Pittsburgh, PA, USA

- Working with [Prof. Hoda Heidari](#) as a Research Assistant in the Machine Learning Department (MLD) at CMU.
- My responsibilities entail of collecting, assimilating, and analyzing both qualitative and quantitative data from prior academic publications, with the goal of creating a tool that offers a pipeline-aware view of Fairness for Machine Learning to researchers and practitioners.

Visiting Research Scholar

Syracuse University

Jun 2022 - Aug 2022

Syracuse, NY, USA

- Working with [Prof. Ferdinando Fioretto](#) on topics related to Differential Privacy and Fairness in AI.

Research Engineer

OpenMined

Mar 2020 - August 2021

Remote

- Collaborated on projects involving Differential Privacy, Federated Learning and Privacy-Preserving Machine Learning protocols with various seniors professionals on PPML at OpenMined.

Application Engineering Intern

BlackRock

Jan 2021 - Jul 2021

Gurugram, India (Remote)

- Part of the Client-End Fund Reporting Team. Improved test coverage on FRED (Factsheet Reporting Engine and Distribution) and fixed code issues, blockers and bugs.
- Received an honourable mention for our internal hackathon project on “BlackRock’s Cultural Heatmap” which provides a forum for both employees (to assess their mental and cultural well-being) and managers (to maintain a cultural pulse throughout the organization).

R&D Intern

Procreate Techno Systems

Jun 2018 - Jul 2018

Hyderabad, India

- Demonstrated an alternative approach for solving Speaker Identification (implemented the K-Means Clustering algorithm). The former approach used Gaussian Mixture Models(GMMs) for Speaker Identification.

- Mentored young students in a Socratic Dialogue competition.
- Evangelized research-based learning to parents and kids. I was part of the Marketing team where I had to attend fairs to interact with prospective students/parents. We posed the classic Tower of Hanoi problem to the kids and asked them to arrange the blocks under a specified time limit.

PUBLICATIONS & PROJECTS

Can Causal (or Counterfactual) Representations benefit from Quantum Computing? [Preprint](#)

We discuss on how latent causal representations can possibly be modelled through quantum computing. (*Accepted as an extended abstract at [AFCI-NeurIPS'22](#)*)

Pruning has a disparate impact on model accuracy [Preprint](#)

We show that accuracy disparities in pruned models arise due to the presence of two key factors: (1) disparity in gradient norms across groups, and (2) disparity in Hessian matrices associated with the loss function computed using a group's data. (*Accepted at [NeurIPS'22](#). 🌟 Spotlight Lightning Talk*)

Fair Context-Aware Privacy Threat Modelling [Preprint](#)

We examine notions of fairness in privacy threat modelling due to different causes of privacy threats within a particular situation/context and that across contexts. (*Presented at [PTM](#) workshop at [USENIX-SOUPS'22](#)*)

Can Causal (and Counterfactual) Reasoning improve Privacy Threat Modelling? [Preprint](#)

We discuss what causal and counterfactual reasoning is and how this can be applied in the field of privacy threat modelling (PTM). (*Presented at [PTM](#) workshop at [USENIX-SOUPS'22](#)*)

Efficient Hyperparameter Optimization for Differentially Private Deep Learning [Preprint](#)

We study three different hyperparameter optimization approaches for DP-SGD to achieve the best privacy-utility tradeoffs. (*Accepted at [PPML](#) workshop at [ACM CCS'21](#)*)

Privacy Enabled Financial Text Classification using Differential Privacy and Federated Learning

We apply DP and FL on Financial Text data and provide results and intuition for the same. (*Accepted at [ECONLP](#) workshop at [EMNLP'21](#)*)

Benchmarking Differential Privacy and Federated Learning for BERT models [Preprint](#)

We benchmark BERT-based models with DP and FL on two Twitter datasets (Depression and Sexual harassments). We provide open source implementations for the same to accelerate Private NLP research. (*Accepted at [ML4Data](#) workshop at [ICML'21](#)*)

Towards Quantifying Carbon Emissions of Differentially Private Machine Learning [Preprint](#)

We quantify carbon emissions for DP-SGD in three different environments : NLP (News Classification), CV (MNIST Digit Classification) and RL (Cartpole problem). (*Accepted at [SRML](#) workshop at [ICML'21](#)*)

DP-SGD vs PATE: Which Has Less Disparate Impact on Model Accuracy? [Preprint](#)

We compare DP-SGD with PATE, another DP-based approach in terms of Fairness. We infer that as PATE uses a teacher-student setup where disjoint data is distributed among the teachers, it suffers less disparity than DP-SGD (due to diversity in training). (*Accepted at [ML4Data](#) workshop at [ICML'21](#) and [PPML](#) workshop at [ACM CCS'21](#)*)

FedPerf: A Practitioners' Guide to Performance of Federated Learning Algorithms [Publication](#)

I worked with this team on extending their NeurIPS short paper with the Stragglers and Robustness experiments. We propose an empirical investigation on four prominent FL algorithms to discover the relation between the FL System Parameters (FLSPs) and their performance. (*Accepted for publication at [PMLR](#)*)

When Differential Privacy Meets Interpretability: A Case Study

[Preprint](#) — [Poster](#)

We investigate the tradeoffs between Differentially-Private SGD (DP-SGD) and Interpretability specifically through CAMs on the APTOS dataset. (*Accepted as extended abstract at [RCV](#) workshop at CVPR'21; full paper accepted at [PPML](#) workshop at ACM CCS'21*)

Improved variants of Score-CAM via Smoothing and Integrating

[Poster](#)

We improve Score-CAM by adding smoothing and integration functions as suggested in the SmoothGrad and IntegratedGrad papers respectively. (*Accepted as extended abstract at [RCV](#) workshop at CVPR'21*)

FedPandemic: A Cross-Device Federated Learning Approach Towards Elementary Prognosis of Diseases During a Pandemic

[Preprint](#)

We come up with a simple noise algorithm (inspired by Randomized Response and integrated with Federated Learning) to retrieve prominent COVID-19 symptoms in a privacy-preserving fashion.

(*Accepted at [DPML](#) and [MLPCP](#) workshops at ICLR'21*)

SS-CAM: Smoothed Score-CAM for sharper visual feature localization

[Preprint](#)

We introduce Smoothing to the Score-CAM algorithm, which is a state-of-the-art CAM algorithm. Smoothing allows us to capture more features of the focused object in the image, which leads to better visually attributed results.

IS-CAM: Integrated Score-CAM for axiomatic-based explanations

[Preprint](#)

We borrow the idea of integration from “IntegratedGrad” and combine it with Score-CAM to conduct faithfulness evaluations. IS-CAM performs better than SS-CAM and Score-CAM in terms of faithfulness evaluations, considering the VGG-16 as our baseline model.

TeleVital: Enhancing the quality of contactless health assessment

[Paper](#) — [News](#)

Our team came 2nd in a pan-Indian hackathon called #CODE19 and won \$5000 for this solution to detect vitals from the webcam itself, thereby promoting remote diagnosis during COVID-19. I worked on the Respiratory rate calculations via webcam and was responsible for documenting the entire project for presenting at the hackathon.

Shamir.jl An implementation of Shamir's Secret Sharing protocol in Julia.

[Link](#)

Shamir's Secret Sharing scheme allows t out of n parties to recover the secret using Polynomial functions. If there are k parties (where $k < t$) in the protocol, they cannot recover the secret.

DiffPrivacy.jl A library to implement Differential Privacy techniques over statistical databases in Julia (Under Development).

[Link](#)

Differential privacy allows on to publicly share information about a dataset by describing the patterns of groups within the dataset while preserving information about individuals in the dataset.

PROFESSIONAL SERVICES

- Teaching Assistant [Quantum Computing Theory and Lab \(11-860\)](#), [Programming Quantum Computers \(17617-A1\)](#), [Quantum Circuit Mappings \(17-620\)](#)
- Served on the Program Committee at the [GenLaw](#) workshop @ ICML'23
- Reviewer at [NeurIPS'23](#)
- Reviewer at the [Algorithmic Fairness through the Lens of Causality and Privacy](#) workshop at NeurIPS'22
- Talk at Comcast Cybersecurity team (headquartered in Philadelphia, PA) on “Context-Aware Privacy Threat Modeling”.
- Served on the Program Committee as a reviewer at [PPAI-AAAI'22](#), [PPAI-AAAI'23](#).
- TEDxMAHE Countdown 2020 Speaker on *Federated Learning for Climate Change*. [Event Link](#) — [Talk](#)
- Manipal Conclave 2020 Student Speaker on *Privacy for ML*. [Memento](#)

- Poster Presented at PyCon India 2019 on *Secure and Private AI with PySyft*.
Volunteered at PyCon India 2020.

Poster

EXTRA-CURRICULAR ACTIVITIES

- Played for the CMU Badminton team in the Fall 2022 Eastern Collegiate MidAtlantic Conf (Badminton Tournament Regionals) held at University of Maryland, College Park in October 2022.
- Finished a full marathon (42 km) at Manipal Marathon 2020 with a timing of 6 hours and 33 minutes. [Certificate](#)
- Core Committee Member of the Linux Users' Group, Manipal (LUGM).
- Open Source Contributor at Python Pandas, OpenMined's PySyft, Julia.
- Captain of the Badminton Team and Libero of the Volleyball team in high school. Came 1st in 2018, 2nd in 2019 at the Inter-branch Badminton Tournament in college.
- Responsible for Field Work and Community Service activities in my high school for a week every term. Field Work involves planting saplings and shovelling while Community Service includes tidying up the school corridors and the school auditorium.